

# Machine Learning Assignment 1

Sayan Sinha  
16CS10048

## INTRODUCTION

Supervised machine Learning generally boils down to approximating a continuous function. One method of doing so is by making use of Linear Regression. In any method whatsoever, there always remains a set of hyperparameters which need to be tuned intuitively. In this assignment, we experiment with various values of such hyperparameters and try to draw a conclusion out of them.

## EXPERIMENTATION

### Part 1

A synthetic dataset is generated and is divided into train and test in a 4:1 ratio randomly. This is done by permuting the data in a random fashion and choosing the first 80% of the dataset to be the training set and the rest to be the test set. This part is kept as a separate function to make sure new data is generated only when required, and not all the time before training. The data is saved in the present working directory.

The dataset is then trained using Linear regression. The weights are initialised to

random values between zero and one. At the end of each iteration, the loss up to that iteration is sent to STDOUT. Gradient descent is performed as per the following rule:

$$\theta_j := \theta_j - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

The training function returns the train error, test error and the RMSE error. For the error, we use squared loss, i.e.:

$$J(\theta) = \frac{1}{2m} \sum_{i=1}^m (W^T \Phi_n(x) - y)^2$$

The training takes place separately for degrees 1 to 9. The degrees denote the polynomial degree of the function we are trying to approximate. Since bias is also introduced, there are a total of  $degree + 1$  parameters to be tuned while training. At the end of a certain number of iterations, these values are saved in the local directory.

The parameter values learned have been provided in Table 1.

| 1                | 2                | 3                | 4                | 5                | 6                | 7                | 8                | 9                |
|------------------|------------------|------------------|------------------|------------------|------------------|------------------|------------------|------------------|
| 0.797531<br>9619 | 0.823028<br>0506 | 0.818393<br>5063 | 1.166862<br>8    | 1.442892<br>006  | 1.594227<br>469  | 1.293419<br>923  | 1.216725<br>803  | 1.056935<br>601  |
| -1.441639<br>135 | -1.628034<br>357 | -1.265472<br>884 | -1.637685<br>999 | -1.876832<br>283 | -2.252152<br>497 | -1.266530<br>609 | -1.278813<br>704 | -0.940861<br>88  |
|                  | 0.197333<br>2989 | -1.059530<br>691 | -1.869621<br>587 | -2.140028<br>412 | -1.565307<br>6   | -1.752213<br>111 | -1.222106<br>24  | -1.410866<br>293 |

|  |  |                  |                  |                   |                    |                    |                   |                    |
|--|--|------------------|------------------|-------------------|--------------------|--------------------|-------------------|--------------------|
|  |  | 0.969801<br>3993 | 0.454302<br>5178 | -0.177134<br>7154 | -0.651195<br>9638  | -1.541795<br>004   | -1.481039<br>603  | -0.929530<br>0843  |
|  |  |                  | 1.535320<br>926  | 0.413476<br>0332  | -0.039120<br>77531 | -0.036397<br>05869 | -0.428483<br>5354 | -0.410347<br>7217  |
|  |  |                  |                  | 2.233017<br>764   | 1.128940<br>674    | 0.862934<br>4587   | 0.357140<br>9422  | -0.016717<br>29166 |
|  |  |                  |                  |                   | 1.847131<br>128    | 0.816115<br>8519   | 0.716485<br>5357  | 0.103302<br>1222   |
|  |  |                  |                  |                   |                    | 1.831891<br>609    | 0.845241<br>2498  | 0.523226<br>7054   |
|  |  |                  |                  |                   |                    |                    | 1.552610<br>11    | 0.886524<br>7169   |
|  |  |                  |                  |                   |                    |                    |                   | 1.451567<br>311    |

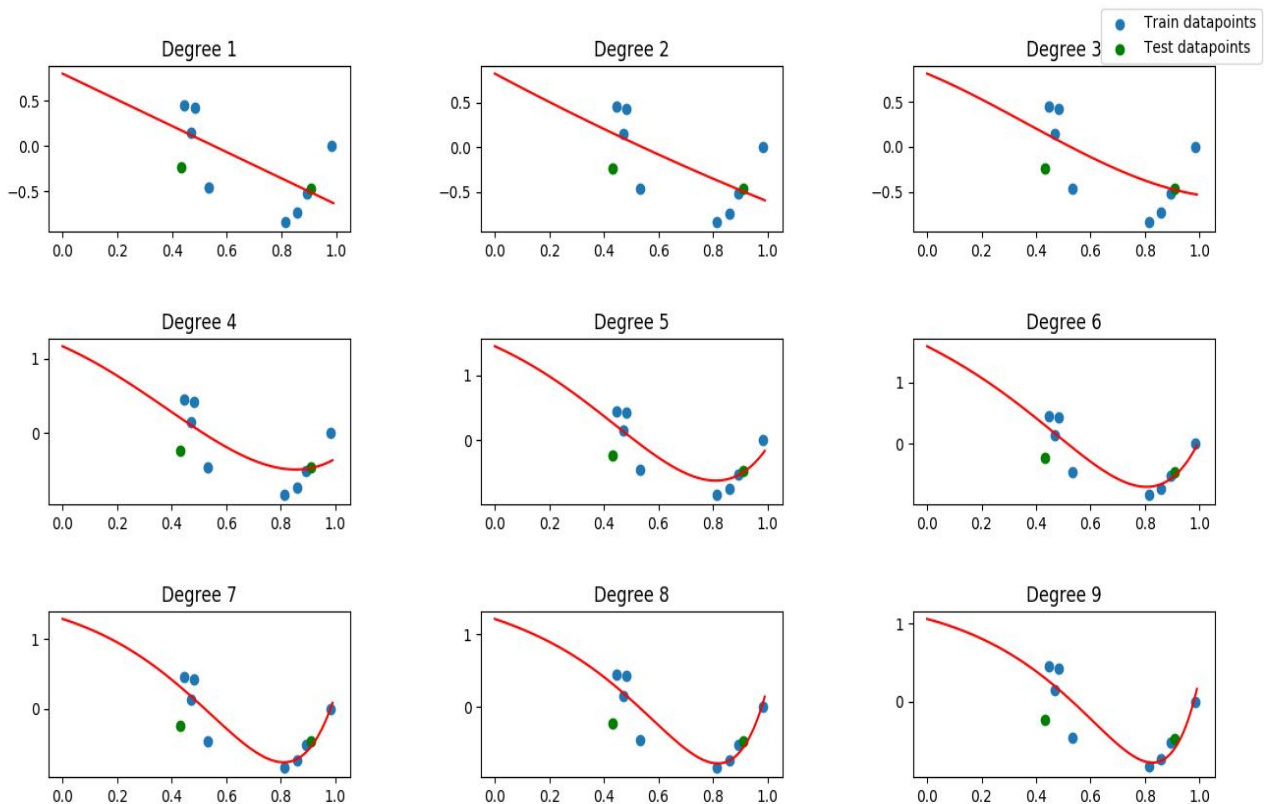
Table 1: Weights for 10 datapoints

## Part 2

Part 2 takes up the job of looking into how good the training went about using visual plots. First, it asks the user if there is a desire to generate a new dataset, or if at all the training needs to be performed again. After such confirmations, it proceeds to plot the data. First, it reads the saved data and weights

after which it recomputes the test and train errors. This is done as the test and train errors, though computed, were not saved in Part 1. For plotting the approximated function using the given weights, the output of the function is obtained for inputs between 0 and 1, separated by 0.01. Thus, when a line plot is drawn on this, it gives the illusion of a curve.

Fig 1: Curve fitting against 10 datapoints



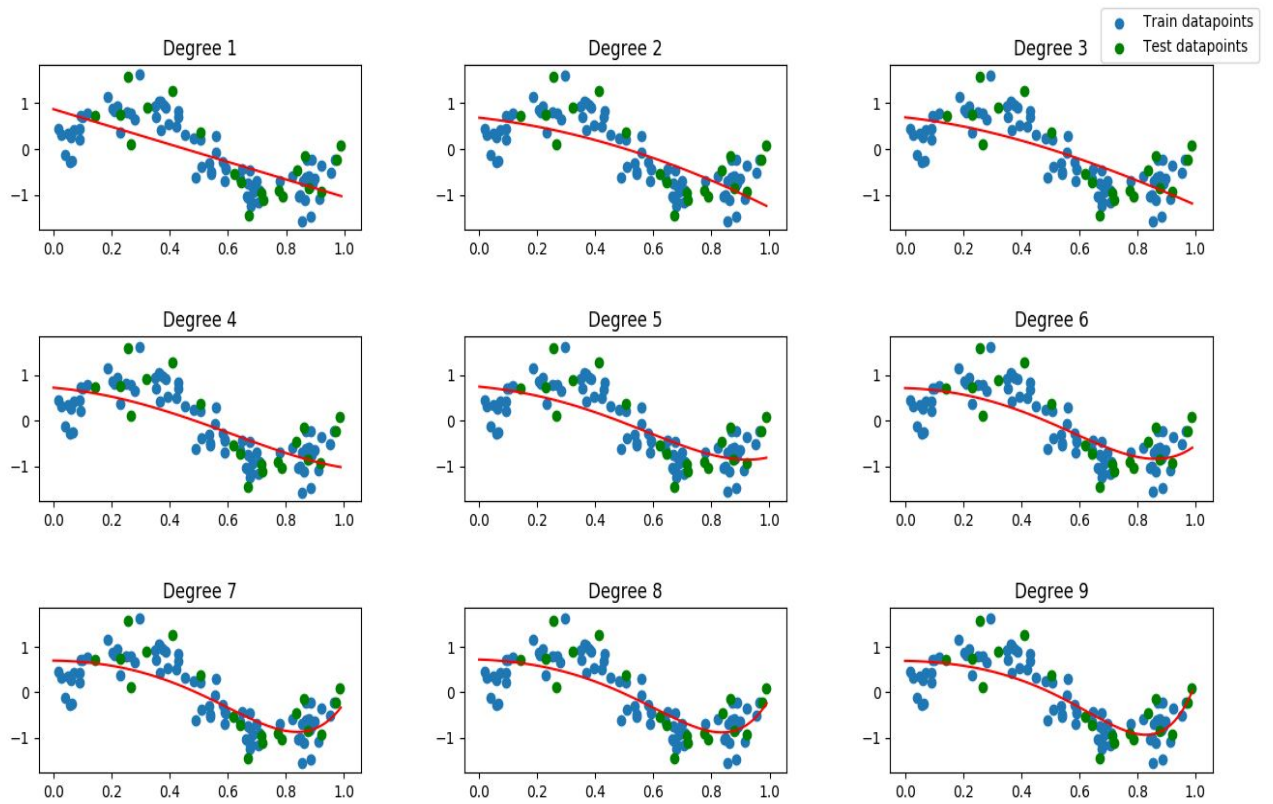


Fig 2: Curve fitting against 100 datapoints




Fig 3: Loss v/s degree for 100 datapoints




Fig 4: Loss v/s degree for 10 datapoints

## Part 3

Part 3 makes a call to Part 2 to obtain the best value of  $n$ . Accordingly, it then instructs Part 1 to generate data with various sizes, namely 10, 100, 1,000, and 10,000. The weights obtained have been presented in Table 2, 3 and 4. Upon testing a number of times, the value of the best  $n$  is found according to the following distribution (Table 5). The plot against no of datapoints is given in Fig 5.

Table 5: Distribution of degrees

|               |      |       |       |
|---------------|------|-------|-------|
| <b>Degree</b> | 7    | 8     | 9     |
| <b>%</b>      | 6.67 | 13.33 | 80.00 |

Table 2: Weights for 100 datapoints

[illegible]

Table 3: Weights for 1000 datapoints

| 1                | 2                 | 3                 | 4                 | 5                 | 6                 | 7                 | 8                  | 9                 |
|------------------|-------------------|-------------------|-------------------|-------------------|-------------------|-------------------|--------------------|-------------------|
| 0.795294<br>9234 | 0.795661<br>4451  | 0.671008<br>6979  | 0.674772<br>2721  | 0.676215<br>1944  | 0.741085<br>8763  | 0.680414<br>0972  | 0.797758<br>8119   | 0.701040<br>2116  |
| -1.606634<br>293 | -1.160043<br>219  | -0.635297<br>9473 | -0.771208<br>8956 | -0.597161<br>0846 | -1.063300<br>577  | -0.659153<br>834  | -1.048135<br>427   | -0.675897<br>5782 |
|                  | -0.643251<br>0116 | -0.593078<br>5415 | -0.377599<br>4579 | -1.008405<br>992  | -0.543273<br>5428 | -1.003218<br>382  | -0.616417<br>6367  | -0.770524<br>1842 |
|                  |                   | -0.603542<br>5272 | -0.709381<br>2278 | -0.529990<br>1328 | -0.156100<br>4362 | -0.325332<br>3324 | -1.015222<br>288   | -0.702825<br>1569 |
|                  |                   |                   | 0.082220<br>34768 | 0.234952<br>8062  | -0.111481<br>7494 | -0.363662<br>6983 | -0.038812<br>81459 | -0.778274<br>6306 |
|                  |                   |                   |                   | 0.267195<br>1549  | -0.173285<br>7535 | 0.475429<br>1399  | 0.417499<br>5428   | 0.278040<br>7732  |
|                  |                   |                   |                   |                   | 0.421244<br>2658  | 0.043486<br>57395 | 0.413433<br>1526   | 0.087793<br>55429 |
|                  |                   |                   |                   |                   |                   | 0.387009<br>0428  | -0.177975<br>0626  | 0.338968<br>5656  |
|                  |                   |                   |                   |                   |                   |                   | 0.767720<br>4145   | 0.901786<br>2445  |
|                  |                   |                   |                   |                   |                   |                   |                    | 0.197253<br>533   |

Table 4: Weights for 10,000 datapoints

| 1                | 2                 | 3                  | 4                 | 5                 | 6                  | 7                  | 8                 | 9                 |
|------------------|-------------------|--------------------|-------------------|-------------------|--------------------|--------------------|-------------------|-------------------|
| 0.781714<br>3592 | 0.709914<br>3363  | 0.682912<br>0663   | 0.690717<br>5004  | 0.682811<br>1931  | 0.752605<br>908    | 0.731046<br>685    | 0.709935<br>9958  | 0.853432<br>7123  |
| -1.572001<br>909 | -0.728133<br>1692 | -0.597355<br>4095  | -0.741943<br>3226 | -0.936088<br>8669 | -0.907285<br>8963  | -0.647606<br>6363  | -0.827493<br>4965 | -1.121219<br>198  |
|                  | -1.033013<br>008  | -1.085452<br>251   | -0.708731<br>8272 | -0.319905<br>1655 | -1.038105<br>144   | -1.134167<br>163   | -0.696408<br>7892 | -0.978014<br>8136 |
|                  |                   | -0.077089<br>84932 | -0.557043<br>3458 | -0.274101<br>1187 | -0.078618<br>73709 | -0.830478<br>5904  | -0.628142<br>2084 | -0.171826<br>8267 |
|                  |                   |                    | 0.292768<br>2211  | -0.307388<br>3862 | -0.296437<br>6787  | 0.032630<br>71575  | -0.196983<br>0868 | -0.722522<br>0256 |
|                  |                   |                    |                   | 0.135145<br>2717  | 0.453599<br>0327   | -0.065458<br>67159 | 0.078122<br>1145  | 0.038532<br>14125 |
|                  |                   |                    |                   |                   | 0.355022<br>3538   | 0.560727<br>4715   | 0.171170<br>8567  | 0.550561<br>6662  |
|                  |                   |                    |                   |                   |                    | 0.817546<br>4423   | 0.118368<br>9984  | -0.161832<br>919  |
|                  |                   |                    |                   |                   |                    |                    | 0.691812<br>0453  | 0.653741<br>6599  |

|  |  |  |  |  |  |  |  |                  |
|--|--|--|--|--|--|--|--|------------------|
|  |  |  |  |  |  |  |  | 0.828182<br>0676 |
|--|--|--|--|--|--|--|--|------------------|

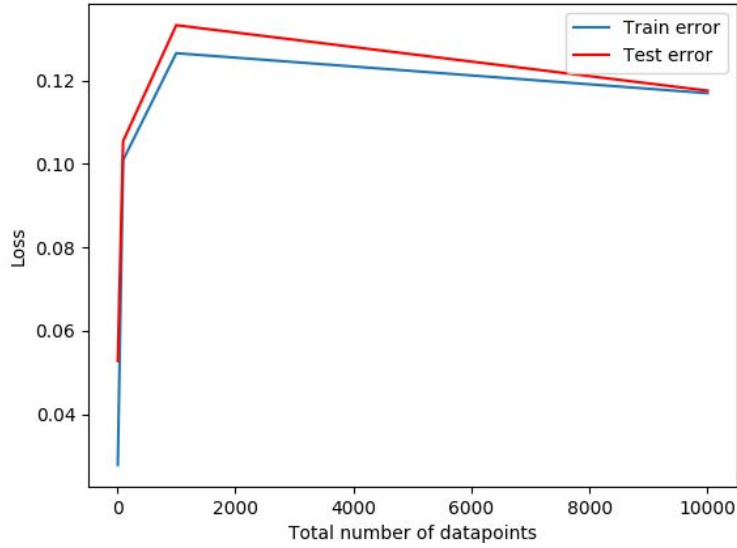


Fig 5: Loss v/s no of datapoints

#### Part 4

Two other cost functions were experimented with, namely the mean absolute error and the fourth power error. For the mean absolute error, the following formula was used for computing cost

$$J(\theta) = \frac{1}{2m} \sum_{i=1}^m |W^T \Phi_n(x) - y|$$

and the following for gradient.

$$\frac{dMAE}{dy_{pred}} = \begin{cases} +1, & y_{pred} > y_{true} \\ -1, & y_{pred} < y_{true} \end{cases}$$

For the fourth power error, the cost function was

$$\frac{1}{2m} \sum_{i=1}^m (W^T \Phi_n(x) - y)^4$$

Accordingly, the gradient comes out to be

$$\frac{2}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^3 x_j^{(i)}$$

Here, the weights have been initialised randomly between 0 and 0.001 to prevent exploding of values.

The RMSE error is trained and computed for each degree. For each instance of the learning rate and for every cost function, the degree having the minimum value of RMSE is chosen and plotted across the learning rate values. The weights have been given in Table 6. The provided weights are for that degree for which the minimum value of RMSE was obtained.

Table 6: Weights for Squared error loss

| 0.025         | 0.05          | 0.1           | 0.2           | 0.5           |
|---------------|---------------|---------------|---------------|---------------|
| 0.7293215757  | 0.7465717494  | 0.9291357642  | 1.042771469   | 0.9614016857  |
| -0.7186918233 | -0.8293871128 | -0.9888155837 | -1.183347055  | -0.5719301969 |
| -0.845257154  | -0.9279066871 | -1.329807527  | -1.330827547  | -1.885688898  |
| -0.8664141269 | -0.2307502415 | -0.8151101417 | -0.9837553255 | -1.699878445  |
| 0.07024936875 | -0.1309270971 | -0.3318826047 | -0.6089554393 | -0.173358147  |
| 0.1790967975  | -0.1041283308 | 0.246028049   | -0.1855961717 | -0.2156919053 |
| -0.1095506467 | 0.1267188376  | 0.4721587248  | 0.9466496538  | 0.669988285   |
| 0.3165117195  | -0.1687833015 | 0.7737945748  | 0.4701290394  | 0.8501394341  |
| 0.2331473181  | 0.7872001551  | 0.8259655405  | 1.190339921   | 1.096247705   |
| 0.5834135234  | 0.242288971   |               | 0.6975961542  | 1.155297139   |

Table 7: Weights for Mean Absolute Error

| 0.025         | 0.05          | 0.1           | 0.2           | 0.5           |
|---------------|---------------|---------------|---------------|---------------|
| 0.5897437497  | 0.7458329931  | 0.990000795   | 1.225462428   | 1.269331133   |
| -0.5366108515 | -0.7737905044 | -1.352076637  | -1.661870657  | -1.43829365   |
| -0.5703084878 | -0.5708972985 | -0.8887757569 | -1.341224788  | -1.558887571  |
| -0.5431326869 | -1.033364742  | -0.3741909892 | -0.4470348273 | -1.569393699  |
| -0.3934620506 | 0.06160323211 | -0.5582285952 | -0.1297491422 | -0.2512687934 |
| 0.263836122   | 0.266539504   | 0.1825713155  | -0.3301895356 | 0.3245040276  |
| 0.2240216534  | 0.06892601893 | 0.4764327226  | 0.2615325467  | 0.9356641203  |
| 0.02726837111 | 0.2371244755  | 0.3778189537  | 0.3538152723  | 1.070499122   |
| -0.1784790498 | 0.5112638983  | 0.3292100977  | 1.07589597    | 1.193997085   |
| 0.5924522313  |               | 0.5246380648  | 0.9210709979  |               |

Table 8: Weights for Fourth Power Error

| 0.025         | 0.05          | 0.1           | 0.2           | 0.5           |
|---------------|---------------|---------------|---------------|---------------|
| 0.6245150067  | 0.8007088396  | 0.9432650773  | 0.989307045   | 0.9290307589  |
| -0.670324757  | -0.9542326425 | -1.168411553  | -1.129097379  | -0.6344867589 |
| -0.6319956396 | -0.883328324  | -1.152236085  | -1.400483947  | -1.812908766  |
| -0.3972964747 | -0.5198589783 | -0.6711720606 | -0.870859511  | -1.294542826  |
| -0.1873818257 | -0.1955922497 | -0.2225846011 | -0.3046772061 | -0.5224040152 |
| -0.0342709929 |               |               |               |               |
| 7             | 0.04018260693 | 0.1083178591  | 0.1296636942  | 0.1047213058  |
| 0.0724038637  | 0.201124684   | 0.3334643396  | 0.4263289758  | 0.5362645949  |
| 0.1431212527  | 0.3063118265  | 0.4794926944  | 0.618876233   | 0.8057831371  |

|              |              |              |              |              |
|--------------|--------------|--------------|--------------|--------------|
| 0.1888810998 | 0.3731522693 | 0.5704175346 | 0.7354502344 | 0.9603547974 |
| 0.2180880341 | 0.4133952611 | 0.6234962827 | 0.8010226496 | 1.034530432  |

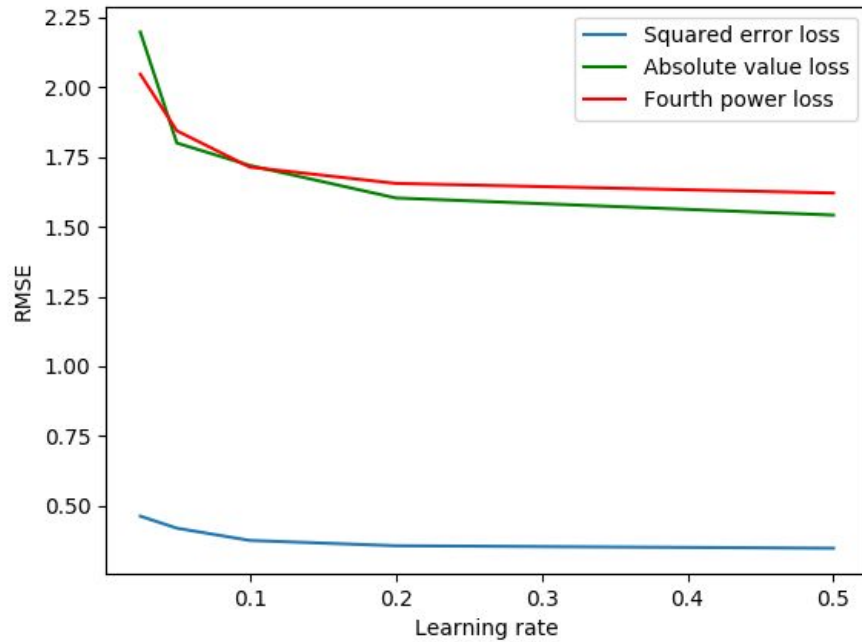


Fig 6: Test RMSE v/s Learning rate

## CONCLUSION

We conclude that the best degree for fitting the curve has been 9 (as evident from Table 5), as seen in Part 4. The best cost function is Squared Error Loss (as evident from Fig 6). The best learning rate has been experimentally found to be 0.5. Along with that, we can infer that we should have sufficient number of datapoints to draw valid conclusions about our model. For instance, in Part 2, having 10 datapoints gave us a wrong notion about our model (ref to Fig 4), whereas having a 100 datapoints could fit the data better (Fig 2). Having an even higher no of datapoints decreases the loss, as evident in Part 3 (Fig 5).

## REFERENCES

- Seber, George AF, and Alan J. Lee. Linear regression analysis. Vol. 329. John Wiley & Sons, 2012.
- Needell, Deanna, Rachel Ward, and Nati Srebro. "Stochastic gradient descent, weighted sampling, and the randomized Kaczmarz algorithm." Advances in Neural Information Processing Systems. 2014.
- Robert, Christian. "Machine learning, a probabilistic perspective." (2014): 62-63.